

finalProject

December 14, 2022

```
[26]: import warnings
warnings.filterwarnings('ignore')

import pandas as pd
import numpy as np
import math
from plotnine import *

from sklearn import metrics
from sklearn.preprocessing import StandardScaler #Z-score variables

# For logistic regression
from sklearn.linear_model import LogisticRegression, Ridge, Lasso # Logistic
    ↳Regression Model
from sklearn.tree import DecisionTreeClassifier # Decision Tree
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split

from sklearn.model_selection import train_test_split # simple TT split cv
from sklearn.model_selection import KFold # k-fold cv
from sklearn.model_selection import cross_val_score # cross validation metrics
from sklearn.model_selection import cross_val_predict # cross validation metrics
from sklearn.model_selection import GridSearchCV
from sklearn.metrics._plot.roc_curve import plot_roc_curve

# Ridge/Lasso Models
from sklearn.linear_model import RidgeCV, LassoCV

# Clustering Model, Gaussian Mixture
from sklearn.mixture import GaussianMixture

from sklearn.metrics import accuracy_score, confusion_matrix, roc_auc_score,
    ↳plot_confusion_matrix, f1_score, recall_score, precision_score
```

[27]:

```

fields = ['GenderSelect', 'Country', 'Age', 'EmploymentStatus', 'CodeWriter',
↪ 'StudentStatus', 'CurrentJobTitleSelect', 'LanguageRecommendationSelect',
↪ 'LearningDataScienceTime', 'TimeSpentStudying', 'FormalEducation',
↪ 'CompensationAmount', 'JobHuntTime', 'EmployerSearchMethod',
↪ 'WorkToolsSelect']
df = pd.read_csv("Datasets/multipleChoiceResponses.csv", usecols = fields ,
↪ encoding = 'latin-1')

df.head()

```

```

[27]:
      GenderSelect      Country      Age \
0  Non-binary, genderqueer, or gender non-conforming      NaN      NaN
1                Female  United States  30.0
2                Male      Canada  28.0
3                Male  United States  56.0
4                Male      Taiwan  38.0

      EmploymentStatus  StudentStatus  CodeWriter \
0      Employed full-time      NaN      Yes
1      Not employed, but looking for work      NaN      NaN
2      Not employed, but looking for work      NaN      NaN
3  Independent contractor, freelancer, or self-em...      NaN      Yes
4      Employed full-time      NaN      Yes

      CurrentJobTitleSelect  LanguageRecommendationSelect \
0      DBA/Database Engineer      F#
1                NaN      Python
2                NaN      R
3  Operations Research Practitioner      Python
4      Computer Scientist      Python

      LearningDataScienceTime  TimeSpentStudying      FormalEducation \
0                NaN                NaN  Bachelor's degree
1      1-2 years      2 - 10 hours  Master's degree
2      1-2 years      2 - 10 hours  Master's degree
3                NaN                NaN  Master's degree
4                NaN                NaN  Doctoral degree

      EmployerSearchMethod \
0  I visited the company's Web site and found a j...
1                NaN
2                NaN
3                NaN
4      A tech-specific job board

      WorkToolsSelect  CompensationAmount \
0  Amazon Web services,Oracle Data Mining/ Oracle...      NaN

```

1		NaN	NaN
2		NaN	NaN
3	Amazon Machine Learning,Amazon Web services,Cl...		250,000
4	C/C++,Jupyter notebooks,MATLAB/Octave,Python,R...		NaN

	JobHuntTime
0	NaN
1	NaN
2	1-2
3	NaN
4	NaN

```
[28]: df.shape
```

```
[28]: (16716, 15)
```

```
[29]: # Cleaning up the data

# Simplify data in GenderSelect
df = df[df['GenderSelect'].isin(['Male', 'Female'])]

# Only United States Participants
df = df[df['Country'] == 'United States']

# Drop those who didnt input their age
df = df.dropna(subset=['Age'])

# Drop 'prefer not to answer' for FormalEducation
df = df[~df['FormalEducation'].isin(['I prefer not to answer'])]

# Drop 'Some other way' for EmployerSearchMethod
df = df[~df['EmployerSearchMethod'].isin(['Some other way'])]

# Simplify EmploymentStatus, if employed: 1, else: 0
df['isEmployed'] = df['EmploymentStatus'].apply(lambda x: 1 if any(s in x for s in ['full-time', 'freelancer']) else 0)

# Convert StudentStatus null values to 0
df['StudentStatus'] = df['StudentStatus'].astype(str).apply(lambda x: 1 if 'Yes' in x else 0)

# Only show those who have a job in analyzing data, software or programming
jobTypes = ['Data', 'Software', 'Computer', 'Database', 'Business Analyst', 'Machine Learning', 'Programmer']
df['CodeWriter'] = df['CodeWriter'].apply(lambda x: 1 if x == 'Yes' else 0)
df = df[df['CurrentJobTitleSelect'].isin(jobTypes) | df['CodeWriter'] == 1]
```

```

# Drop those who didnt put their annual salary
df = df.dropna(subset=['CompensationAmount'])

# Convert the salary from a string to an integer value
df['CompensationAmount'] = df['CompensationAmount'].apply(lambda x: int(float(x.
    ↪replace(',', ''))))

```

```

[30]: # Simplifying labels
df['FormalEducation'] = df['FormalEducation'].replace('I did not complete any,
    ↪formal education past high school', 'high_school')
df['FormalEducation'] = df['FormalEducation'].replace('Some college/university,
    ↪study without earning a bachelor\'s degree', 'some_college')
df['FormalEducation'] = df['FormalEducation'].replace('Professional degree',
    ↪'professional')
df['FormalEducation'] = df['FormalEducation'].replace('Bachelor\'s degree',
    ↪'bachelors')
df['FormalEducation'] = df['FormalEducation'].replace('Master\'s degree',
    ↪'masters')
df['FormalEducation'] = df['FormalEducation'].replace('Doctoral degree', 'PhD')

df['EmployerSearchMethod'] = df['EmployerSearchMethod'].replace('A career fair,
    ↪or on-campus recruiting event', 'campus_recruitment')
df['EmployerSearchMethod'] = df['EmployerSearchMethod'].replace('A friend,
    ↪family member, or former colleague told me', 'referral')
df['EmployerSearchMethod'] = df['EmployerSearchMethod'].replace(['A,
    ↪general-purpose job board', 'A tech-specific job board', 'I visited the,
    ↪company\'s Web site and found a job listing there'], 'job_board')
df['EmployerSearchMethod'] = df['EmployerSearchMethod'].replace(['I was,
    ↪contacted directly by someone at the company (e.g. internal recruiter)', 'An,
    ↪external recruiter or headhunter'], 'recruiter')

```

```

[31]: # Create dummies for GenderSelect, FormalEducation and EmployerSearchMethod
dummies = pd.get_dummies(df['GenderSelect'])
df = pd.concat([df, dummies], axis = 1)
df = df.drop('GenderSelect', 1)

dummies = pd.get_dummies(df['FormalEducation'])
df = pd.concat([df, dummies], axis = 1)
df = df.drop('FormalEducation', 1)

dummies = pd.get_dummies(df['EmployerSearchMethod'])
df = pd.concat([df, dummies], axis = 1)
df = df.drop('EmployerSearchMethod', 1)

```

```

[32]: df.head()

```

```

[32]:      Country    Age      EmploymentStatus \
3   United States  56.0  Independent contractor, freelancer, or self-em...
15  United States  58.0  Independent contractor, freelancer, or self-em...
22  United States  33.0      Employed full-time
34  United States  35.0      Employed full-time
75  United States  40.0      Employed full-time

      StudentStatus  CodeWriter      CurrentJobTitleSelect \
3           0           1  Operations Research Practitioner
15          0           1      DBA/Database Engineer
22          0           1      Scientist/Researcher
34          0           1      Engineer
75          0           1      Scientist/Researcher

      LanguageRecommendationSelect  LearningDataScienceTime  TimeSpentStudying \
3           Python      NaN      NaN
15          R      NaN      NaN
22         Matlab      NaN      NaN
34         Python      NaN      NaN
75         Python      NaN      NaN

      WorkToolsSelect ...  PhD bachelors \
3   Amazon Machine Learning,Amazon Web services,Cl... ...  0      0
15  C/C++,IBM Cognos,MATLAB/Octave,Microsoft Excel... ...  0      0
22         MATLAB/Octave,Python ...  1      0
34  MATLAB/Octave,Python,R,SAS JMP,SQL,TIBCO Spotfire ...  1      0
75         Amazon Machine Learning,C/C++,NoSQL,R ...  1      0

      high_school  masters  professional  some_college  campus_recruitment \
3           0      1      0      0      0
15          0      1      0      0      0
22          0      0      0      0      0
34          0      0      0      0      0
75          0      0      0      0      0

      job_board  recruiter  referral
3           0      0      0
15          0      0      0
22          0      0      1
34          0      0      1
75          0      0      1

[5 rows x 25 columns]

```

```

[33]: predictors = ['Male','Female','Age', 'CompensationAmount',
    'PhD', 'bachelors', 'high_school', 'masters', 'professional',
    'some_college', 'campus_recruitment', 'job_board', 'recruiter',

```

```

    'referral']
continuous_variables = ['Age', 'CompensationAmount']

X = df[predictors]
y = df['isEmployed']

```

```

[34]: #80/10 split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)

# zscore (only continous and interval variables)
z = StandardScaler()
X_train[continuous_variables] = z.fit_transform(X_train[continuous_variables])
X_test[continuous_variables] = z.transform(X_test[continuous_variables])

```

```

[35]: # Logistic Regression Model
lr = LogisticRegression()
lr.fit(X_train, y_train) # test set should never see the inside of a .fit

```

```

[35]: LogisticRegression()

```

```

[36]: # Performance Metrics
print("Train Acc: ", accuracy_score(y_train, lr.predict(X_train)))
print("Test Acc: ", accuracy_score(y_test, lr.predict(X_test)))

print("TRAIN Precision: ", precision_score(y_train, lr.predict(X_train)))
print("TEST Precision : ", precision_score(y_test, lr.predict(X_test)))

print("TRAIN Recall: ", recall_score(y_train, lr.predict(X_train)))
print("TEST Recall : ", recall_score(y_test, lr.predict(X_test)))

print("TRAIN ROC/AUC: ", roc_auc_score(y_train, lr.predict_proba(X_train)[:,:1]))
print("TEST ROC/AUC : ", roc_auc_score(y_test, lr.predict_proba(X_test)[:,:1]))

```

```

Train Acc:  0.9689849624060151
Test Acc:  0.951310861423221
TRAIN Precision:  0.9689849624060151
TEST Precision :  0.951310861423221
TRAIN Recall:  1.0
TEST Recall :  1.0
TRAIN ROC/AUC:  0.7835875731123064
TEST ROC/AUC :  0.6025136281041793

```

```


[37]: # ROC Curve for LR
plot_roc_curve(lr, X_train, y_train)

```

```


[37]: <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x106cc0b80>

```



```
[38]: # ROC Curve for LR
      plot_roc_curve(lr, X_test, y_test)
```

```
[38]: <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x2887d45b0>
```



```
[39]: lasso_model = Lasso()
lasso_model.fit(X_train, y_train)

print("TRAIN: ", mean_absolute_error(y_train, lasso_model.predict(X_train)))
print("TEST : ", mean_absolute_error(y_test, lasso_model.predict(X_test)))
```

```
TRAIN:  0.060106210074057295
TEST :  0.07668398524400886
```

```
[40]: ridge_model = Ridge()
ridge_model.fit(X_train, y_train)

print("TRAIN: ", mean_absolute_error(y_train, ridge_model.predict(X_train)))
print("TEST : ", mean_absolute_error(y_test, ridge_model.predict(X_test)))
```

```
TRAIN:  0.06380915063680255
TEST :  0.081029587436048
```

```
[41]: # LASSO
predictors = ['Male', 'Female', 'Age', 'CompensationAmount',
              'PhD', 'bachelors', 'high_school', 'masters', 'professional',
```



```

    'some_college', 'campus_recruitment', 'job_board', 'recruiter',
    'referral']
continuous_variables = ['Age', 'CompensationAmount']

X = df[predictors]
y = df['isEmployed']

#80/10 split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)

# zscore (only continous and interval variables)
z = StandardScaler()
X_train[continuous_variables] = z.fit_transform(X_train[continuous_variables])
X_test[continuous_variables] = z.transform(X_test[continuous_variables])

lsr_tune = LassoCV(cv = 5).fit(X_train,y_train)
# lsr_tune = LassoCV(cv = 5, alphas = [0.001,0.01,0.05,1]).fit(X_train,y_train)

print("TRAIN: ", mean_absolute_error(y_train, lsr_tune.predict(X_train)))
print("TEST : ", mean_absolute_error(y_test, lsr_tune.predict(X_test)))

print("\nwe chose " + str(lsr_tune.alpha_) + " as our alpha.")

```

TRAIN: 0.07477043489143065

TEST : 124.41368238267599

we chose 0.0010694794140082694 as our alpha.

```

[42]: # Ridge tuning
predictors = ['Male','Female','Age', 'CompensationAmount',
    'PhD', 'bachelors', 'high_school', 'masters', 'professional',
    'some_college', 'campus_recruitment', 'job_board', 'recruiter',
    'referral']
continuous_variables = ['Age', 'CompensationAmount']

X = df[predictors]
y = df['isEmployed']

#80/10 split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)

# zscore (only continous and interval variables)
z = StandardScaler()
X_train[continuous_variables] = z.fit_transform(X_train[continuous_variables])
X_test[continuous_variables] = z.transform(X_test[continuous_variables])

rr_tune = RidgeCV(cv = 5).fit(X_train,y_train)

```

```
print("TRAIN: ", mean_absolute_error(y_train, rr_tune.predict(X_train)))
print("TEST : ", mean_absolute_error(y_test, rr_tune.predict(X_test)))

print("\nwe chose " + str(rr_tune.alpha_) + " as our alpha.")
```

TRAIN: 0.0582307193038051
TEST : 0.08546631957736815

we chose 10.0 as our alpha.

```
[43]: #from scipy.stats import zscore
#z_scores = zscore(df['CompensationAmount'])
#df_clean = df[(z_scores < 3).all(axis=1) & (z_scores > -3).all(axis=1)]

# Remove Outliers
df = df[df['CompensationAmount'] < 250_000]
df = df[df['CompensationAmount'] > 10_000]

df = df[df['Age'] > 18]
```

```
[44]: (ggplot(df, aes(x = "Age", y = "CompensationAmount"))
      + geom_point()
      + theme_minimal()
      + ggtitle("Relation between Age and Annual Salary")
      + labs(x = "Age", y = "Annual Salary")
      )
```



[44]: <ggplot: (679856171)>

```
[45]: from sklearn.preprocessing import StandardScaler

features = ['Age', 'CompensationAmount']
z = StandardScaler()

df[features] = z.fit_transform(df[features])
gmm = GaussianMixture(n_components = 3).fit(df[features])
gmm_labels = gmm.predict(df[features])

#labList = ["Cluster " + str(i) for i in range(1, len(set(gmm.labels_)))]

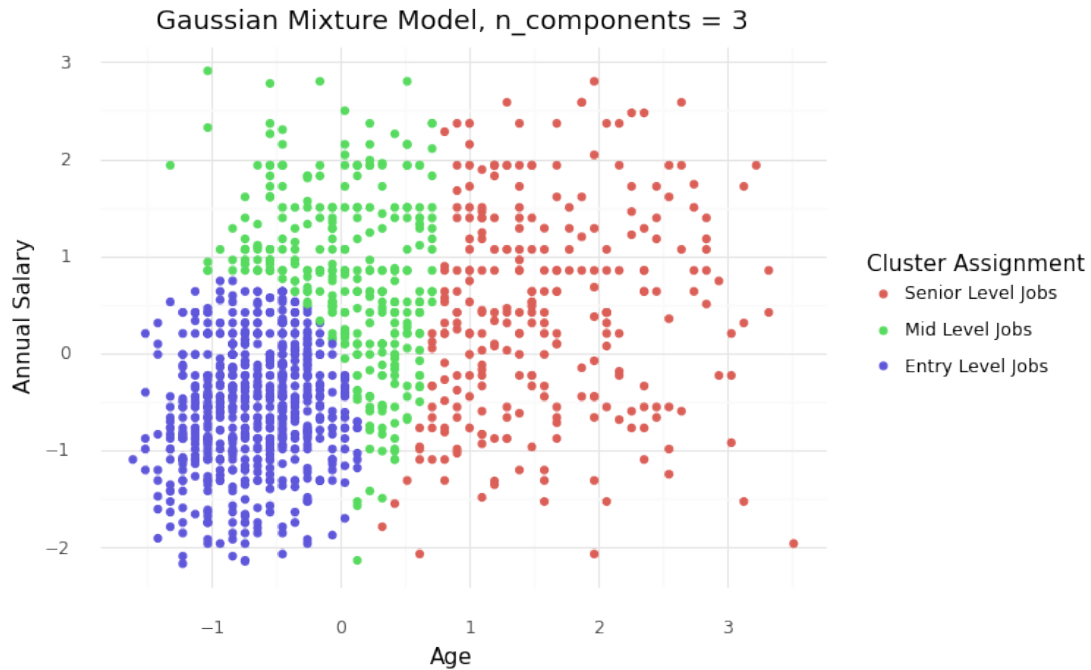
df['assignments'] = gmm_labels
```

```
[49]: (ggplot(df, aes(x = "Age", y = "CompensationAmount", color =_
↪ "factor(assignments)"))
      + geom_point()
      + theme_minimal())
```

```

+ labs(title = "Gaussian Mixture Model, n_components = 3", x = "Age", y = "Annual Salary")
+ scale_color_discrete(name = "Cluster Assignment", labels = (["Senior Level Jobs", "Mid Level Jobs", "Entry Level Jobs"]))
)

```



```
[49]: <ggplot: (680356255)>
```

```

[47]: ss_gmm = silhouette_score(df[features], df[['assignments']])
print("Silhouette Score for GM Model:", ss_gmm)

```

Silhouette Score for GM Model: 0.34963370059537546